



# Deep Dive into Kubernetes

## Part 2

Imesh Gunaratne, WSO2

# Agenda

- Docker Image Size Optimization
- Docker Image Size Optimization Demo
- Persistent Volumes
  - Static Volumes
  - Dynamic Volumes
- Container Security
  - File System Permissions
  - Pod Security Context
  - Pod Security Policies



# Agenda Cont.

- Kubernetes on Google Cloud Platform
- GKE Demo
- Kubernetes on AWS
- Kubernetes on Azure



# Docker Image Size Optimization

# A Sample Dockerfile

```
FROM anapsix/alpine-java:8

ARG USER_GROUP_ID=1000
ARG USER_GROUP=bar
ARG USER_ID=2000
ARG USER=foo

RUN addgroup -S -g ${USER_GROUP_ID} ${USER_GROUP} && \
    adduser -S -g ${USER_GROUP_ID} -u ${USER_ID} ${USER}

COPY --chown=foo:bar helloworld-* /tmp/

USER ${USER_ID}

ENTRYPOINT [ "java", "-jar", "/tmp/helloworld-2.5.3-SNAPSHOT.jar" ]
```



# Docker Image Size Optimization

1. Use a smaller base image
2. Install only application dependent software
3. Minimize layers and combine RUN commands

```
RUN groupadd --system -g ${USER_GROUP_ID} ${USER_GROUP} && \  
    useradd --system --create-home --home-dir ${USER_HOME} --no-log-init  
-g ${USER_GROUP_ID} -u ${USER_ID} ${USER}
```

4. List layers according to change frequency



# Docker Image Size Optimization Cont.

5. Use `--no-install-recommends` on `apt-get install`:

```
RUN apt-get update && \  
    apt-get install -y --no-install-recommends \  
    <package-name>
```

6. Add `rm -rf /var/lib/apt/lists/*` to same layer as `apt-get install`:

```
RUN apt-get update && \  
    apt-get install -y --no-install-recommends \  
    <package-name> \  
    rm -rf /var/lib/apt/lists/*
```



# Docker Image Size Optimization Cont.

7. Copy archive files after extracting them locally:

```
COPY ${FILES}/${JDK} ${USER_HOME}/java/
```

8. Use --chown key instead of chown command when copying files:

```
COPY --chown=wso2carbon:wso2 ${FILES}/${JDK} ${USER_HOME}/java/
```





# Docker Image Size Optimization Cont.

Finally use docker history command to verify the image structure:

```
docker history gcr.io/google-samples/node-hello:1.0
```

IMAGE	CREATED	CREATED BY	SIZE
4c7ea8709739	2 years ago	/bin/sh -c #(nop) CMD ["/bin/sh" "-c" "node ...	0B
<missing>	2 years ago	/bin/sh -c #(nop) COPY file:0ee96426ce9ede6e...	861B
<missing>	2 years ago	/bin/sh -c #(nop) EXPOSE 8080/tcp	0B
<missing>	2 years ago	/bin/sh -c #(nop) CMD ["node"]	0B
<missing>	2 years ago	/bin/sh -c curl -sLO "https://nodejs.org/dis...	37MB
<missing>	2 years ago	/bin/sh -c #(nop) ENV NODE_VERSION=4.4.2	0B
<missing>	2 years ago	/bin/sh -c #(nop) ENV NPM_CONFIG_LOGLEVEL=in...	0B
<missing>	2 years ago	/bin/sh -c set -ex && for key in 9554F...	60.4kB
<missing>	2 years ago	/bin/sh -c apt-get update && apt-get install...	315MB
<missing>	2 years ago	/bin/sh -c apt-get update && apt-get install...	123MB
<missing>	2 years ago	/bin/sh -c apt-get update && apt-get install...	44.3MB
<missing>	2 years ago	/bin/sh -c #(nop) CMD ["/bin/bash"]	0B
<missing>	2 years ago	/bin/sh -c #(nop) ADD file:b5391cb13172fb513...	125MB



# Docker Image Size Optimization Demo

# Persistent Volumes

# Persistent Volume Usage Design

Pod

Persistent  
Volume Claim

Persistent  
Volume

Physical  
Storage



# Persistent Volume Provisioning Types

- **Static**
  - Manually map physical storage to PVs
- **Dynamic**
  - When static PVs are unavailable, K8S will check the availability of dynamic PVs
  - It dynamically creates volumes using storage provisioners
  - The storage provisioner may use a single physical storage for providing multiple dynamic PVs

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>

<https://github.com/kubernetes-incubator/external-storage/>



# Persistent Volume Claims

```
kind: PersistentVolumeClaim
```

```
apiVersion: v1
```

```
metadata:
```

```
  name: myclaim
```

```
spec:
```

```
  accessModes:
```

```
    - ReadWriteOnce
```

```
  volumeMode: Filesystem
```

```
  resources:
```

```
    requests:
```

```
      storage: 8Gi
```

```
  storageClassName: slow
```

```
  selector:
```

```
    matchLabels:
```

```
      release: "stable"
```


```
    matchExpressions:
```

```
      - {key: environment, operator: In, values: [dev]}
```



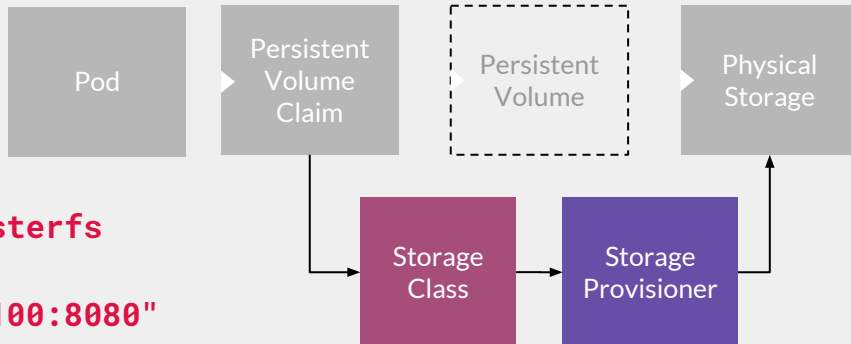
# Static Persistent Volumes

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
  mountOptions:
    - hard
    - nfsvers=4.1
  nfs:
    path: /tmp
    server: 172.17.0.2
```



# Dynamic Persistent Volumes

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: gluster-vol-default
provisioner: kubernetes.io/glusterfs
parameters:
  resturl: "http://192.168.10.100:8080"
  restuser: ""
  secretNamespace: ""
  secretName: ""
allowVolumeExpansion: true
```





# Persistent Volume Volume Modes

- Filesystem
- Block

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: block-pv
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  volumeMode: Block
  persistentVolumeReclaimPolicy: Retain
  fc:
    targetWWNs: ["50060e801049cfd1"]
    lun: 0
    readOnly: false
```



# Persistent Volume Access Modes

- **ReadWriteOnce**
  - read-write by a single node
- **ReadOnlyMany**
  - read-only by many nodes
- **ReadWriteMany**
  - read-write by many nodes



# Persistent Volume Binding

- Binding a PVC to PV will be done based on requested:
  - Amount of storage
  - Access mode
  - Storage Class (Optional)



# Persistent Volume Re-claim Policies

- **Retain**
  - Manual reclamation
- **Recycle**
  - Basic scrub (`rm -rf /thevolume/*`) on the volume and makes it available again for a new claim
- **Delete**
  - Deletion removes both the PV object from Kubernetes, as well as the associated storage asset in the external infrastructure

<https://kubernetes.io/docs/concepts/storage/persistent-volumes/>



# Container Security

# Container Security

- Container host operating system security
- Container operating system security
- Application security
- Container registry permissions
- Network isolation
- Container filesystem permissions
- Persistent volume permissions
- Container/container cluster manager API endpoint permissions



# File System Permission Management

# File System Permissions

1. Create an OS user and a group in the container image:

```
RUN groupadd --system -g ${USER_GROUP_ID} ${USER_GROUP} && \  
    useradd --system --create-home --home-dir ${USER_HOME} --no-log-init -g  
    ${USER_GROUP_ID} -u ${USER_ID} ${USER}
```

2. Specify user and the group for granting permissions to the filesystem:

```
COPY --chown=user:group ${FILES}/${JDK} ${USER_HOME}/java/
```





# File System Permissions Cont.

3. Define the user id in the container image before the entrypoint:

```
USER ${USER_ID}  
...  
ENTRYPOINT ${SERVER_HOME}/bin/server.sh
```



# Pod Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo
spec:
  securityContext:
    runAsUser: 1000
    fsGroup: 2000
  volumes:
  - name: sec-ctx-vol
    emptyDir: {}
  containers:
  - name: sec-ctx-demo
    image: gcr.io/google-samples/node-hello:1.0
    volumeMounts:
    - name: sec-ctx-vol
      mountPath: /data/demo
    securityContext:
      allowPrivilegeEscalation: false
```

Define the user id and the group id in the pod under security context:

<https://kubernetes.io/docs/tasks/configure-pod-container/security-context/>



# Persistent Volume Permissions

- By default PVs would get root permissions (root:root).
- If the container is designed to run using a non-root user, that user may not have permissions for accessing the PVs.
- Therefore, the user group (gid) used in the container may need to be used for granting access to PVs.
- This can be done by:
  - Defining the gid in the PV.
  - By using a Pod Security Policy and defining the gid using the fsGroup property.



# Define Group ID in Persistent Volume

Define the group id specified in the container image in the PV:

```
kind: PersistentVolume
apiVersion: v1
metadata:
  name: pv1
  annotations:
    pv.beta.kubernetes.io/gid: "1234"
```

<https://kubernetes.io/docs/tasks/configure-pod-container/configure-persistent-volume-storage/>



# Pod Security Policies

## Control Aspect

Running of privileged containers

Usage of the root namespaces

Usage of host networking and ports

Usage of volume types

Usage of the host filesystem

White list of FlexVolume drivers

Allocating an FSGroup that owns the pod's volumes

## Field Names

[privileged](#)

[hostPID, hostIPC](#)

[hostNetwork, hostPorts](#)

[volumes](#)

[allowedHostPaths](#)

[allowedFlexVolumes](#)

[fsGroup](#)



# Pod Security Policies Cont.

## Control Aspect

Requiring the use of a read only root file system

The user and group IDs of the container

Restricting escalation to root privileges

Linux capabilities

## Field Names

[readOnlyRootFilesystem](#)

[runAsUser](#),  
[supplementalGroups](#)

[allowPrivilegeEscalation](#),  
[defaultAllowPrivilegeEscalation](#)

[defaultAddCapabilities](#),  
[requiredDropCapabilities](#),  
[allowedCapabilities](#)



# Pod Security Policies Cont.

## Control Aspect

The SELinux context of the container

The AppArmor profile used by containers

The seccomp profile used by containers

The sysctl profile used by containers

## Field Names

[seLinux](#)

[annotations](#)

[annotations](#)

[annotations](#)



# Defining File System Group (fsGroup)

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: restricted
  annotations:
    ...
spec:
  ...
  fsGroup:
    rule: 'MustRunAs'
    ranges:
      # Allow following gid range for accessing PVs
      - min: 1000
        max: 1000
```





# Defining Run As User, Supplemental Groups

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: example
spec:
  ...
  # start container using
  following uid and gid range:
  runAsUser:
    rule: MustRunAs
    ranges:
  supplementalGroups:
    rule: RunAsAny
  ...
```

- **runAsUser**
  - MustRunAs (uid range)
  - MustRunAsNonRoot
  - RunAsAny
- **supplementalGroups:**
  - MustRunAs (gid range)
  - RunAsAny





# Kubernetes on Google Cloud Platform

# Kubernetes on Google Cloud



- It's called Google Kubernetes Engine (GKE)
- Kubernetes is provided as a service
  - Masters are provided for free
- IAM: Role based access with Google Accounts
- Private container registries
- Autoscaling
- Auto upgrade Kubernetes clusters
- Logging and Monitoring with Stackdriver

<https://cloud.google.com/kubernetes-engine/>



# Kubernetes on Google Cloud Cont.



- **Databases:**
  - Cloud SQL: MySQL, PostgreSQL
- **Persistent Volumes:**
  - **GCEPersistentDisk:** ReadWriteOnce, ReadOnlyMany
  - **NFS, Glusterfs:** ReadWriteMany
- **Load Balancers:**
  - Ingresses/Load Balancer Type Services creates Google Cloud Load Balancers

<https://cloud.google.com/kubernetes-engine/>



# GKE Demo



# Kubernetes on AWS

# Kubernetes on AWS



- Recently announced: Amazon Elastic Container Service (EKS)
- Kubernetes will be provided as a service with EKS
  - Masters will be provided for free
- AWS Auth for role based access
- Currently, KOPs can be used for creating a self managed Kubernetes clusters on AWS



# Kubernetes on AWS



- **Databases:**
  - RDS: Amazon Aurora, MySQL, PostgreSQL, MariaDB, Oracle, MS SQL
- **Persistent Volumes:**
  - EBS: ReadWriteOnce
  - EFS, NFS, Glusterfs -> ReadWriteMany
- **Load Balancers:**
  - Ingresses/Load Balancer Type Services creates AWS Load Balancers







Kubernetes on Azure

# Kubernetes on Azure



- Azure Container Service (AKS)
- Currently in preview mode
- A managed Kubernetes service:
  - Masters will be provided for free
- Private container registries
- Autoscaling
- Automated Kubernetes cluster upgrades

<https://azure.microsoft.com/en-us/services/container-service/>



# Kubernetes on Azure



- **Databases:**
  - Azure SQL: MySQL, PostgreSQL, MS SQL
- **Persistent Volumes:**
  - **AzureDisk**-> ReadWriteOnce
  - **AzureFile, NFS, Glusterfs:** ReadWriteOnce, ReadOnlyMany, ReadWriteMany
- **Load Balancers:**
  - Ingresses/Load Balancer Type Services creates Azure Load Balancers



Questions & Feedback

# References

# References

- Tips to Reduce Docker Image Sizes:
  - <https://hackernoon.com/tips-to-reduce-docker-image-sizes-876095da3b34>
- Introduction to Container Security:
  - [https://www.docker.com/sites/default/files/WP\\_IntrotoContainerSecurity\\_08.19.2016.pdf](https://www.docker.com/sites/default/files/WP_IntrotoContainerSecurity_08.19.2016.pdf)
- Ten Layers of Container Security:
  - <https://www.redhat.com/cms/managed-files/cl-container-security-openshift-cloud-devops-tech-detail-f7530kc-201705-en.pdf>



# References Cont.

- Kubernetes Persistent Volumes:
  - <https://kubernetes.io/docs/concepts/storage/persistent-volumes/>
- Kubernetes External Storage Plugins, Provisioners, and Helper Libraries:
  - <https://github.com/kubernetes-incubator/external-storage/>
- Pod Security Policies:
  - <https://kubernetes.io/docs/concepts/policy/pod-security-policy/>



Thank You!